

Map-Reduce Algorithm

The pseudo code for the different map reduce patterns is as follows

No-Combiner Approach:

Method map (Key k, Value v):

 Extract StationId, RecordType, temperature from v

 If (recordType = "TMAX" or recordType = "TMIN"):

 emit (stationId, (recordType, temperature,1);

Method reduce (Key k, values [(recordType1, temperature1,1),.....]):

//key is the station-id

// value consists of record-type , temperature reading and count

 maxTempCount \leftarrow 0

 maxTempSum \leftarrow 0

 minTempCount \leftarrow 0

 minTempSum \leftarrow 0

 minAverage \leftarrow 0

 maxAverage \leftarrow 0

 for each record v in values:

 if (v.recordType = "TMAX"):

 maxTempSum \leftarrow maxTempSum + v.temperature

 maxTempCount \leftarrow maxTempCount + 1

 if (v.recordType = "TMIN"):

 minTempSum \leftarrow minTempSum + v.temperature

 minTempCount \leftarrow minTempCount + 1

 minAverage \leftarrow minTempSum/minTempCount

 maxAverage \leftarrow maxTempSum/maxTempCount

 emit(stationId, (minAverage, maxAverage))

Combiner Approach:

Method map (Key k, Value v):

 Extract StationId, RecordType, temperature from v

 If (recordType = "TMAX" or recordType = "TMIN"):

 emit (stationId, (recordType, temperature,1);

end

Method combine ((Key k, values [(recordType1, temperature1,1),.....])):

// key consists of station-id

// value consists of record-type, temperature reading and count

 maxTempCount \leftarrow 0

 maxTempSum \leftarrow 0

 minTempCount \leftarrow 0

 minTempSum \leftarrow 0

 for each record v in values:

 if (v.recordType = "TMAX"):

 maxTempSum \leftarrow maxTempSum + v.temperature

 maxTempCount \leftarrow maxTempCount + 1

 if (v.recordType = "TMIN"):

 minTempSum \leftarrow minTempSum + v.temperature

 minTempCount \leftarrow minTempCount + 1

 emit(stationId, ("TMAX", maxTempSum))

 emit(stationId, ("TMIN", minTempSum))

end

method reduce (Key k, values [(recordType1, temperature1,count1),.....]):

// key consists of station-id

// value consists of record-type, temperature reading and count

maxTempCount \leftarrow 0

maxTempSum \leftarrow 0

minTempCount \leftarrow 0

minTempSum \leftarrow 0

minAverage \leftarrow 0

maxAverage \leftarrow 0

for each record v in values:

if (v.recordType = "TMAX"):

maxTempSum \leftarrow maxTempSum + v.temperature

maxTempCount \leftarrow maxTempCount + v.count

if (v.recordType = "TMIN"):

minTempSum \leftarrow minTempSum + v.temperature

minTempCount \leftarrow minTempCount + v.count

minAverage \leftarrow minTempSum/minTempCount

maxAverage \leftarrow maxTempSum/maxTempCount

emit(stationId, (minAverage, maxAverage))

end

In-Mapper Combining Approach:

```
class map {  
  
    method initialize () :  
        H-max  $\leftarrow$  Initialize HashMap;  
        H-min  $\leftarrow$  Initialize HashMap;  
  
    method map(Key k , Value v):  
        sum  $\leftarrow$  0  
        total  $\leftarrow$  0  
        Extract StationId, RecordType, temperature from v  
        If (recordType = "TMAX"):  
            sum  $\leftarrow$  H-max{stationId}.temperature + temperature  
            total  $\leftarrow$  H-max{stationId}.count + 1  
            H-max.add(stationId,(sum,total))  
  
        If (recordType = "TMIN"):  
            sum  $\leftarrow$  H-min{stationId}.temperature + temperature  
            total  $\leftarrow$  H-min{stationId}.count + 1  
            H-min.add(stationId,(sum,total))  
        end  
  
    method cleanup():  
  
        for each key in H-max :  
            emit(key, (H-max{key}.sum, H-max{key}.total))  
  
        for each key in H-min :  
            emit(key, (H-min{key}.sum, H-min{key}.total))  
        end  
  
}
```

```
class reduce{

method reduce (Key k, values [(recordType1, temperature1,count1),..... ]):
    maxTempCount  $\leftarrow$  0
    maxTempSum  $\leftarrow$  0
    minTempCount  $\leftarrow$  0
    minTempSum  $\leftarrow$  0
    minAverage  $\leftarrow$  0
    maxAverage  $\leftarrow$  0

    for each record v in values:
        if (v.recordType = "TMAX"):
            maxTempSum  $\leftarrow$  maxTempSum + v.temperature
            maxTempCount  $\leftarrow$  maxTempCount + v.count

        if (v.recordType = "TMIN"):
            minTempSum  $\leftarrow$  minTempSum + v.temperature
            minTempCount  $\leftarrow$  minTempCount + v.count

    minAverage  $\leftarrow$  minTempSum/minTempCount
    maxAverage  $\leftarrow$  maxTempSum/maxTempCount
    emit(stationId, (minAverage, maxAverage))
end

}
```

Secondary Sort :

// map function converts the value into a key and value pair . The key is an object of CustomKey class

method map(Key k, value v) :

 from value v extract station-id, year, record-type and temperature

 if(record-type = "TMAX"):

 emit((station-id,year), (year, 0,0,temperature,1))

 if(record-type= "TMIN") :

 emit((station-id,year),(year,temperature,1,0,0))

end

//The key is an object of type CustomKey. Custom key consists of two attributes. stationId and year.

class CustomKey {

 stationId

 year

 method compareTo(Custom key k1 , Custom Key k2)

 compare the stationId's .

 if(same):

 compare year

 end

}

//The partitioner takes the key which is of type CustomKey and returns an appropriate partition based on stationId . All records with a particular station-id go to same reducer

method partitioner(key):

 return hash(key.stationId)

end

//The combiner takes two parameters , key of the type CustomKey and list of values having same key

method combiner(key, values[(year, maxTempSum0,maxTempCount0,
minTempSum0,minTempCount0),...]):

maxTempSum \leftarrow 0
minTempSum \leftarrow 0
maxTempCount \leftarrow 0
minTempCount \leftarrow 0

for each v in values:

maxTempSum \leftarrow maxTempSum + v.maxTempSum
minTempSum \leftarrow minTempSum + v.minTempSum
maxTempCount \leftarrow maxTempCount + v.maxTempCount
minTempCount \leftarrow minTempCount + v.minTempCount

emit(key, (year, maxTempSum, maxTempCount, minTempSum,
minTempCount)

end

// The grouping comparator groups data by station id and all records having same stationId are sent to same reducer.

method customGroupComparator (Key k1, Key k2):

//Key consists of station-id and year
compareValue \leftarrow compare(k1.station-id, k2.station-id)
return compareValue

end

//The reduce function takes two parameters . A key of type CustomKey and list of values having same key . Each value contains five components . year , TMIN sum , TMIN, count , TMAX sum, TMAX count. The values are received in the increasing order of year

```
method reduce(key k, values[(year, minTemp, maxTemp, minTempCount,
                             maxTempCount)]):
```

```
    maxTempSum ← 0
    minTempSum ← 0
    maxTempCount ← 0
    minTempCount ← 0
    year ← key.year
```

```
    for each v in values:
```

```
        if(v.year is not equal to year)
            emit(key, (year, maxTempSum, maxTempCount, minTempSum,
                       minTempCount)
            maxTempSum ← 0
            minTempSum ← 0
            maxTempCount ← 0
            minTempCount ← 0
```

```
    maxTempSum ← maxTempSum + v.maxTempSum
    minTempSum ← minTempSum + v.minTempSum
    maxTempCount ← maxTempCount + v.maxTempCount
    minTempCount ← minTempCount + v.minTempCount
```

```
end
```

The mapper emits records in the increasing order of stationId. This is achieved by overriding the inbuilt compareTo function. We defined our own custom compareTo function in CustomKey class . In this compareTo function , we first compare by stationId . If stationId's are equal, then we compare the years. Thus mapper emits records in increasing order of keys

The grouping comparator the groups all the records having similar stationId and sends them to the reducer. Since we have defined our own custom comparator,

the records in the reducer will be already present in the increasing order of year(In custom comparator , we compare by year when stationId's are equal).

Thus we make use of map reduce's sorting ability to prevent explicit sorting of values in reducer. This in-turn eliminates the need for complex data structures needed to sort.

Performance Comparison

The performance for different design patterns of map-reduce for two executions using 6 m4.large machines is as follows

No Combiner :

Execution 1: Execution 1 completed in 88 seconds

```
INFO total process run time: 88 seconds
2017-02-12T04:12:06.771Z INFO Step created jobs: job_1486872401215_0001
2017-02-12T04:12:06.771Z INFO Step succeeded with exitCode 0 and took 88 seconds
```

Execution 2: Execution 2 completed in 84 seconds

```
INFO total process run time: 84 seconds
2017-02-12T04:43:20.296Z INFO Step created jobs: job_1486874379995_0001
2017-02-12T04:43:20.297Z INFO Step succeeded with exitCode 0 and took 84 seconds
```

Combiner Design Pattern :

Execution 1 : Execution 1 completed in 82 seconds

```
INFO waitProcessCompletion ended with exit code 0 : hadoop jar /mnt/var/lib/hadoop/steps/
s-2YL0FHTW...
INFO total process run time: 98 seconds
2017-02-11T21:03:14.458Z INFO Step created jobs: job_1486846755603_0001
2017-02-11T21:03:14.458Z INFO Step succeeded with exitCode 0 and took 98 seconds
```

Execution 2: Execution 2 completed in 82 seconds

```
INFO total process run time: 82 seconds
2017-02-12T05:36:59.869Z INFO Step created jobs: job_1486877623798_0001
2017-02-12T05:36:59.869Z INFO Step succeeded with exitCode 0 and took 82 seconds
```

In-Mapper Design Pattern

Execution 1: Execution 1 completed in 76 seconds

```
INFO total process run time: 76 seconds
2017-02-11T22:29:34.415Z INFO Step created jobs: job_1486851932347_0001
2017-02-11T22:29:34.415Z INFO Step succeeded with exitCode 0 and took 76 seconds
```

Execution 2: Execution 2 completed in 80 seconds

```
INFO total process run time: 80 seconds
2017-02-11T22:41:51.068Z INFO Step created jobs: job_1486852705270_0001
2017-02-11T22:41:51.068Z INFO Step succeeded with exitCode 0 and took 80 seconds
```

Was the Combiner called at all in program Combiner? Was it called more than once per Map task?

Combiner was called on both the executions in combiner program. The entries in log confirm our observations

Combiner Execution 1:

```
Map-Reduce Framework
  Map input records=30868726
  Map output records=8798241
  Map output bytes=255148989
  Map output materialized bytes=6029332
  Input split bytes=1620
  Combine input records=8798241
  Combine output records=525992
  Reduce input groups=14135
  Reduce shuffle bytes=6029332
  Reduce input records=525992
  Reduce output records=14135
  Spilled Records=1051984
  Shuffled Maps =180
  Failed Shuffles=0
  Merged Map outputs=180
  GC time elapsed (ms)=16585
  CPU time spent (ms)=198280
  Physical memory (bytes) snapshot=17119162368
  Virtual memory (bytes) snapshot=107879714816
  Total committed heap usage (bytes)=15439233024
```

Combiner Execution 2:

```
Map-Reduce Framework
  Map input records=30868726
  Map output records=8798241
  Map output bytes=255148989
  Map output materialized bytes=6029332
  Input split bytes=1620
  Combine input records=8798241
  Combine output records=525992
  Reduce input groups=14135
  Reduce shuffle bytes=6029332
  Reduce input records=525992
  Reduce output records=14135
  Spilled Records=1051984
  Shuffled Maps =180
  Failed Shuffles=0
  Merged Map outputs=180
  GC time elapsed (ms)=16546
  CPU time spent (ms)=191850
  Physical memory (bytes) snapshot=16824655872
  Virtual memory (bytes) snapshot=107836256256
  Total committed heap usage (bytes)=15233187840
```

What difference did the use of a Combiner make in Combiner compared to NoCombiner?

By using the combiner, we can see that the network traffic is greatly reduced. Since combiner combines the intermediate records, we can see from the logs that number of keys moving between mapper and reducer is reduced. These observations can be seen in logs as shown below

Combiner :

```
Map-Reduce Framework
  Map input records=30868726
  Map output records=8798241
  Map output bytes=255148989
  Map output materialized bytes=6029332
  Input split bytes=1620
  Combine input records=8798241
  Combine output records=525992
  Reduce input groups=14135
  Reduce shuffle bytes=6029332
  Reduce input records=525992
  Reduce output records=14135
  Spilled Records=1051984
  Shuffled Maps =180
  Failed Shuffles=0
  Merged Map outputs=180
  GC time elapsed (ms)=16585
  CPU time spent (ms)=198280
  Physical memory (bytes) snapshot=17119162368
  Virtual memory (bytes) snapshot=107879714816
  Total committed heap usage (bytes)=15439233024
```

No Combiner :

```
Map-Reduce Framework
  Map input records=30868726
  Map output records=8798241
  Map output bytes=219956025
  Map output materialized bytes=49005105
  Input split bytes=1620
  Combine input records=0
  Combine output records=0
  Reduce input groups=14135
  Reduce shuffle bytes=49005105
  Reduce input records=8798241
  Reduce output records=14135
  Spilled Records=17596482
  Shuffled Maps =180
  Failed Shuffles=0
  Merged Map outputs=180
  GC time elapsed (ms)=17122
  CPU time spent (ms)=211990
  Physical memory (bytes) snapshot=16907845632
  Virtual memory (bytes) snapshot=107945103360
  Total committed heap usage (bytes)=15397289984
```


Was the local aggregation effective in InMapperComb compared to NoCombiner?

Local aggregation in “In-Mapper combining” provides several benefits

- a) The number of records that are emitted between mapper and reducer is greatly reduced . We can see these observations in the log records shown below

No Combiner :

```
Map-Reduce Framework
  Map input records=30868726
  Map output records=8798241
  Map output bytes=219956025
  Map output materialized bytes=49005105
  Input split bytes=1620
  Combine input records=0
  Combine output records=0
  Reduce input groups=14135
  Reduce shuffle bytes=49005105
  Reduce input records=8798241
  Reduce output records=14135
  Spilled Records=17596482
  Shuffled Maps =180
  Failed Shuffles=0
  Merged Map outputs=180
  GC time elapsed (ms)=17122
  CPU time spent (ms)=211990
  Physical memory (bytes) snapshot=16907845632
  Virtual memory (bytes) snapshot=107945103360
  Total committed heap usage (bytes)=15397289984
```

In-Mapper Combining:

```
Map-Reduce Framework
  Map input records=30868726
  Map output records=223783
  Map output bytes=8056188
  Map output materialized bytes=4018316
  Input split bytes=1598
  Combine input records=0
  Combine output records=0
  Reduce input groups=14135
  Reduce shuffle bytes=4018316
  Reduce input records=223783
  Reduce output records=14135
  Spilled Records=447566
  Shuffled Maps =153
  Failed Shuffles=0
  Merged Map outputs=153
  GC time elapsed (ms)=19103
  CPU time spent (ms)=153200
  Physical memory (bytes) snapshot=15585722368
  Virtual memory (bytes) snapshot=98027868160
  Total committed heap usage (bytes)=14260109312
```

From the above two images we can see a significant difference in “Map output records” field between In mapper combining and no Mapper combining

- b) The time taken to execute in-mapper combining is lesser compared to no-combining. This can be attributed to the fact that less records travel between mapper and reducer. We can see the difference in execution time between the no combiner approach and as follows

In Mapper Combining:

Execution 1: Execution 1 completed in 76 seconds

```
2017-02-11T22:29:34.415Z INFO total process run time: 76 seconds
2017-02-11T22:29:34.415Z INFO Step created jobs: job_1486851932347_0001
2017-02-11T22:29:34.415Z INFO Step succeeded with exitCode 0 and took 76 seconds
```

Execution 2: Execution 2 completed in 80 seconds

```
2017-02-11T22:41:51.068Z INFO total process run time: 80 seconds
2017-02-11T22:41:51.068Z INFO Step created jobs: job_1486852705270_0001
2017-02-11T22:41:51.068Z INFO Step succeeded with exitCode 0 and took 80 seconds
```

No Mapper Combining :

Execution 1: Execution 1 completed in 88 seconds

```
2017-02-12T04:12:06.771Z INFO total process run time: 88 seconds
2017-02-12T04:12:06.771Z INFO Step created jobs: job_1486872401215_0001
2017-02-12T04:12:06.771Z INFO Step succeeded with exitCode 0 and took 88 seconds
```

Execution 2: Execution 2 completed in 84 seconds

```
2017-02-12T04:43:20.296Z INFO total process run time: 84 seconds
2017-02-12T04:43:20.296Z INFO Step created jobs: job_1486874379995_0001
2017-02-12T04:43:20.297Z INFO Step succeeded with exitCode 0 and took 84 seconds
```

From the above two images, we can see that in both the executions, in-mapper combining approach improves the execution time.

Which one is better, Combiner or InMapperComb? Briefly justify your answer.

In Map-reduce, the extent to which efficiency can be increased through local-aggregation depends on size of intermediate key space, distribution of keys and number of key-value pairs that are emitted by each individual map task.

Generally, In-Mapper combiner is better due to two reasons,

- 1) The number of records emitted by map function is much lesser in “in-mapper combining” design pattern than in combiner design pattern. This is due to the fact that in “in-mapper” design pattern, data is aggregated before being emitted. We have observed the same in our experiments. Below is the snapshot of the part of log records. We can see that the number of records emitted by map in “in-mapper combining” is significantly lesser than the number of records emitted by map in combiner approach.

In-Mapper Combining

```
Map-Reduce Framework
  Map input records=30868726
  Map output records=223783
  Map output bytes=8056188
  Map output materialized bytes=4018316
  Input split bytes=1598
  Combine input records=0
  Combine output records=0
  Reduce input groups=14135
  Reduce shuffle bytes=4018316
  Reduce input records=223783
  Reduce output records=14135
  Spilled Records=447566
  Shuffled Maps =153
  Failed Shuffles=0
  Merged Map outputs=153
  GC time elapsed (ms)=19103
  CPU time spent (ms)=153200
  Physical memory (bytes) snapshot=15585722368
  Virtual memory (bytes) snapshot=98027868160
  Total committed heap usage (bytes)=14260109312
```

Combiner Approach

```
Map-Reduce Framework
  Map input records=30868726
  Map output records=8798241
  Map output bytes=255148989
  Map output materialized bytes=6029332
  Input split bytes=1620
  Combine input records=8798241
  Combine output records=525992
  Reduce input groups=14135
  Reduce shuffle bytes=6029332
  Reduce input records=525992
  Reduce output records=14135
  Spilled Records=1051984
  Shuffled Maps =180
  Failed Shuffles=0
  Merged Map outputs=180
  GC time elapsed (ms)=16585
  CPU time spent (ms)=198280
  Physical memory (bytes) snapshot=17119162368
  Virtual memory (bytes) snapshot=107879714816
  Total committed heap usage (bytes)=15439233024
```


- 2) There is no guarantee that the combiner will be executed by the map-reduce framework. Hence it is not a good practice to rely on combiner approach. With “In-mapper combining” technique, we can ensure that combining always happens and hence we can ensure that our job runs efficiently.

However there are certain cons for in-mapper combining approach. They are

- a) This approach breaks the functional programming underpinnings of map-reduce since state is being preserved across multiple input key-value pairs
- b) Preserving state across multiple input instances means that algorithmic behavior may depend on order in which input key-value pairs are encountered. This creates order dependent bugs which are difficult to debug on large datasets.
- c) There is a fundamental scalability bottleneck associated with in-mapper combining data pattern. It critically depends on having sufficient memory to store intermediate results until mapper has completely processed all key-value pairs in an input split

How do the running times and accuracy of these MapReduce programs compare to the sequential implementation of per-station mean temperature? Modify, run, and time the sequential version of your HW1 program on the 1991.csv data. Make sure to change it to measure the end-to-end running time by including the time spent reading the file. Tip: Modify your code to read and process the data line by line (i.e., instead of reading it all into memory). Finally, compare the MapReduce output to the sequential program output to verify and report on its correctness.

The running time of the modified sequence program is 12.575s. The accuracy for both program seems to be correct.



Execution of Program 2

Execution 1: Execution 1 completed in 58 seconds

```
INFO total process run time: 58 seconds
2017-02-12T00:56:57.396Z INFO Step created jobs: job_1486860854922_0001
2017-02-12T00:56:57.396Z INFO Step succeeded with exitCode 0 and took 58 seconds
```

The map reduce execution details for program 2 are as follows:

```
Map-Reduce Framework
  Map input records=7003910
  Map output records=1470207
  Map output bytes=72040143
  Map output materialized bytes=66929
  Input split bytes=940
  Combine input records=1470207
  Combine output records=2417
  Reduce input groups=441
  Reduce shuffle bytes=66929
  Reduce input records=2417
  Reduce output records=441
  Spilled Records=4834
  Shuffled Maps =90
  Failed Shuffles=0
  Merged Map outputs=90
  GC time elapsed (ms)=6827
  CPU time spent (ms)=63910
  Physical memory (bytes) snapshot=9817706496
  Virtual memory (bytes) snapshot=74876817408
  Total committed heap usage (bytes)=8852602880
```