

ASSIGNMENT-1  
AMOGH HUILGOL  
CS6240- SECTION 02

**For each of the versions of your sequential and multithreaded program detailed in B and C, report the minimum, average, and maximum running time observed over the 10 runs. (5 points)**

**SEQUENCE WITHOUT FIB()**

Maximum execution time for 10.0 rounds of program execution is 5.918 seconds  
Minimum execution time for 10.0 rounds of program execution is 2.724 seconds  
Average execution time for 10.0 rounds of program execution is 3.4140000000000006 seconds

**SEQUENCE WITH FIB()**

Maximum execution time for 10.0 rounds of program execution is 16.042 seconds  
Minimum execution time for 10.0 rounds of program execution is 13.772 seconds  
Average execution time for 10.0 rounds of program execution is 14.065800000000001 seconds

**No LOCK WITHOUT FIB()**

Maximum execution time for 10 rounds of program execution is 2.863 seconds  
Minimum execution time for 10 rounds of program execution is 1.419 seconds  
Average execution time for 10 rounds of program execution is 1.6292000000000002 seconds

**NO\_LOCK WITH FIB()**

Maximum execution time for 10 rounds of program execution is 11.809 seconds  
Minimum execution time for 10 rounds of program execution is 11.021 seconds  
Average execution time for 10 rounds of program execution is 11.3602 seconds

**Coarse Lock Without FIB()**

Maximum execution time for 10 rounds of program execution is 2.698 seconds  
Minimum execution time for 10 rounds of program execution is 1.468 seconds  
Average execution time for 10 rounds of program execution is 1.6396000000000002 seconds

**Coarse Lock With FIB()**

Maximum execution time for 10 rounds of program execution is 28.951 seconds  
Minimum execution time for 10 rounds of program execution is 24.61 seconds  
Average execution time for 10 rounds of program execution is 26.6085 seconds

### **Fine Lock Without FIB()**

Maximum execution time for 10.0 rounds of program execution is 2.792 seconds

Minimum execution time for 10.0 rounds of program execution is 1.58 seconds

Average execution time for 10.0 rounds of program execution is 1.8167000000000002 seconds

### **FINE LOCK WITH FIB()**

Maximum execution time for 10.0 rounds of program execution is 13.062 seconds

Minimum execution time for 10.0 rounds of program execution is 10.848 seconds

Average execution time for 10.0 rounds of program execution is 12.046600000000002 seconds

### **No Shared WITHOUT FIB()**

Maximum execution time for 10 rounds of program execution is 4.385 seconds

Minimum execution time for 10 rounds of program execution is 1.431 seconds

Average execution time for 10 rounds of program execution is 1.7981000000000003 seconds

### **NO SHARED WITH FIB()**

Maximum execution time for 10 rounds of program execution is 8.07 seconds

Minimum execution time for 10 rounds of program execution is 5.99 seconds

Average execution time for 10 rounds of program execution is 6.775499999999999 seconds

- 1) Report the number of worker threads used and the speedup of the multithreaded versions based on the corresponding average running times. (5 points)**

The number of worker threads spawned are 4. the speedup obtained is as follows

#### **NO\_LOCK :**

PROGRAM B : 2.09

PROGRAM C : 1.23

#### **COARSE\_LOCK :**

PROGRAM B : 2.08

PROGRAM C : 0.528

#### **FINE\_LOCK :**

PROGRAM B : 1.879

PROGRAM C : 1.16

**NO\_SHARED :**  
PROGRAM B : 1.898  
PROGRAM C : 2.07

**3) Answer the following questions in a brief and concise manner: (4 points each)**

- i) Which program version (SEQ, NO-LOCK, COARSE-LOCK, FINE-LOCK, NO-SHARING) would you normally expect to finish fastest and why? Do the experiments confirm your expectation? If not, try to explain the reasons.**

Answer) NO\_LOCKING will run the fastest as compared to the other versions because

- a) Since there are no locks, we can utilize maximum parallelism and no worker will have to wait for resource
- b) All workers will update the resource at the same time. hence there is a chance for values to be inconsistent
- c) The average values confirm with the expectation

- ii) Which program version (SEQ, NO-LOCK, COARSE-LOCK, FINE-LOCK, NO-SHARING) would you normally expect to finish slowest and why? Do the experiments confirm your expectation? If not, try to explain the reasons.**

Answer) a) Sequential is generally expected to finish the slowest since each step is executed one after the other . In parallel approach , program makes use of multiple cores available in the system. Hence parallel programs usually finish faster.

- b) Yes the experiments confirm with expectation

- iii) Compare the temperature averages returned by each program version. Report if any of them is incorrect or if any of the programs crashed because of concurrent accesses.**

Answer) The effects of parallelism are better felt when input size is large . However we can see from our experiments that the average values for NO\_LOCK mechanism is wrong or inconsistent . This could occur due to

- a) Program crashing because there is lack of synchronization between threads which leads to one thread trying to read a value in Hashmap which is not present . this results in code throwing null pointer exception
- b) Since there is no lock involved , multiple threads may update a data structure based on old record leading to overwriting of one record over other . this leads to inconsistency or wrong results.

- iv) Compare the running times of SEQ and COARSE-LOCK. Try to explain why one is slower than the other. (Make sure to consider the results of both B and C—this might support or refute a possible hypothesis.)**

Answer) Sequential is generally slower than coarse lock mechanism . However we have seen that when delay fib(17) is introduced , sequential program performs faster . the possible reasons could be :

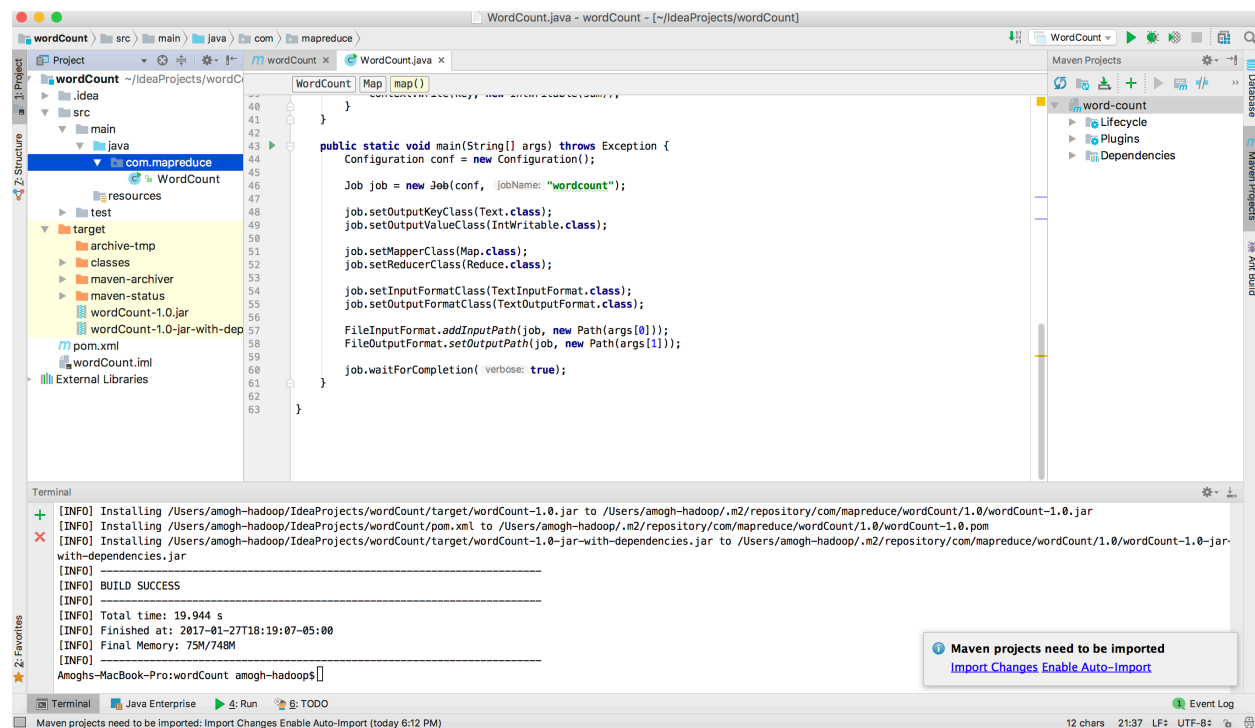
- a) When inputs are small , the tasks are performed quickly with less transfer of locks between threads . Hence parallelism is utilized to speed up the computation
- b) However as the inputs increase , the wait time for threads increases as they have to wait more often for other thread to release locks . This defeats the improvement in performance we obtain by running a program parallel

v) **How does the higher computation cost in part C (additional Fibonacci computation) affect the difference between COARSE-LOCK and FINE-LOCK? Try to explain the reason.**

Answer) In COARSE\_LOCK , the lock is held over the entire data structures and this ensures that no other thread is trying to update the hashmap. thus it can be guaranteed that no two Fibonacci computations will happen in parallel. Thus all delays will happen in sequential way. However for FINE\_LOCK there is a chance that some Fibonacci computations may happen in parallel since each thread holds lock over the variable which it updates (not the whole hashmap) . Hence with higher computation cost , the fine lock performs better than coarse lock

## WORD COUNT LOCAL EXECUTION

Screenshot showing the directory structure



## Screen shot showing last 20 lines of Hadoop execution in IDE

```
2017-01-27 17:30:40,343 INFO [pool-3-thread-1] output.FileOutputCommitter (FileOutputCommitter.java:commitTask(535)) - Saved output of task 'attempt_local1771957700_0001_r_000000_0' to fi
2017-01-27 17:30:40,343 INFO [pool-3-thread-1] mapred.LocalJobRunner (LocalJobRunner.java:statusUpdate(591)) - reduce > reduce
2017-01-27 17:30:40,344 INFO [pool-3-thread-1] mapred.Task (Task.java:sendDone(1158)) - Task 'attempt_local1771957700_0001_r_000000_0' done.
2017-01-27 17:30:40,344 INFO [pool-3-thread-1] mapred.LocalJobRunner (LocalJobRunner.java:run(325)) - Finishing task: attempt_local1771957700_0001_r_000000_0
2017-01-27 17:30:40,345 INFO [Thread-17] mapred.LocalJobRunner (LocalJobRunner.java:runTasks(456)) - reduce task executor complete.
2017-01-27 17:30:40,444 INFO [main] mapreduce.Job (Job.java:monitorAndPrintJob(1367)) - map 100% reduce 100%
2017-01-27 17:30:40,444 INFO [main] mapreduce.Job (Job.java:monitorAndPrintJob(1378)) - Job job_local1771957700_0001 completed successfully
2017-01-27 17:30:40,481 INFO [main] mapreduce.Job (Job.java:monitorAndPrintJob(1385)) - Counters: 30

File System Counters
  FILE: Number of bytes read=109939463586
  FILE: Number of bytes written=141872513878
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0

Map-Reduce Framework
  Map input records=21987700
  Map output records=248943500
  Map output bytes=2418234700
  Map output materialized bytes=2916121964
  Input split bytes=4752
  Combine input records=0
  Combine output records=0
  Reduce input groups=5273
  Reduce shuffle bytes=2916121964
  Reduce input records=248943500
  Reduce output records=5273
  Spilled Records=744919293
  Shuffled Maps =44
  Failed Shuffles=0
  Merged Map outputs=44
  GC time elapsed (ms)=4750
  Total committed heap usage (bytes)=55790534656

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=1454183628
File Output Format Counters
  Bytes Written=73395

Process finished with exit code 0
```

## EXECUTION ON AWS

### Screenshot showing successful execution of word count in AWS-EMR

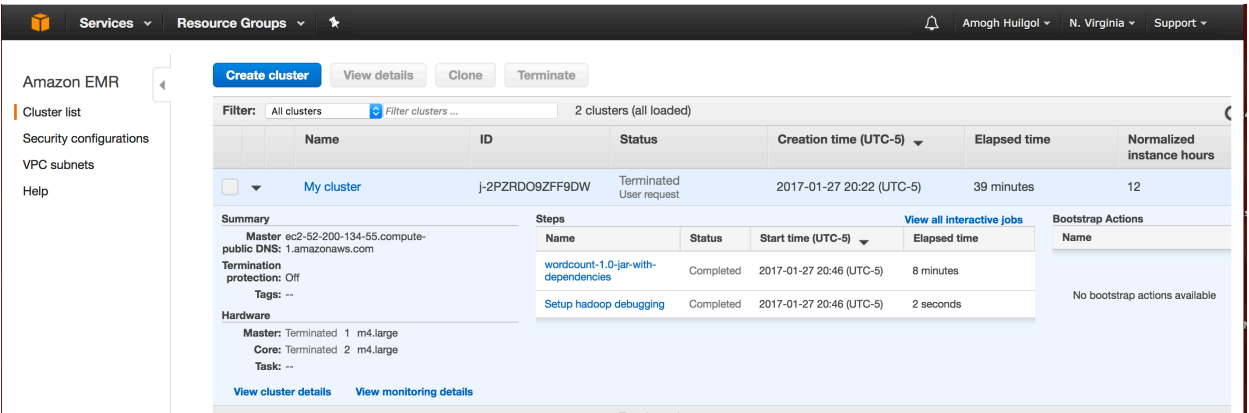
The screenshot displays the AWS Management Console interface for Amazon EMR. A modal window titled "wordcount-1.0-jar-with-dependencies" is open, showing the status of a completed job. The job details include:

- Status:** Completed
- ID:** s-X66L5S32PRVL
- Start time:** 2017-01-27 20:46 (UTC-5)
- Elapsed time:** 8 minutes
- Log files:** controller | syslog | stderr | stdout
- JAR location:** s3://mr-amogh-jars/wordcount-1.0-jar-with-dependencies.jar
- Main class:** None
- Arguments:** s3://mr-inputs/hw12.txt s3://mr-test-inputs/outputhw1
- Action on failure:** Cancel and wait

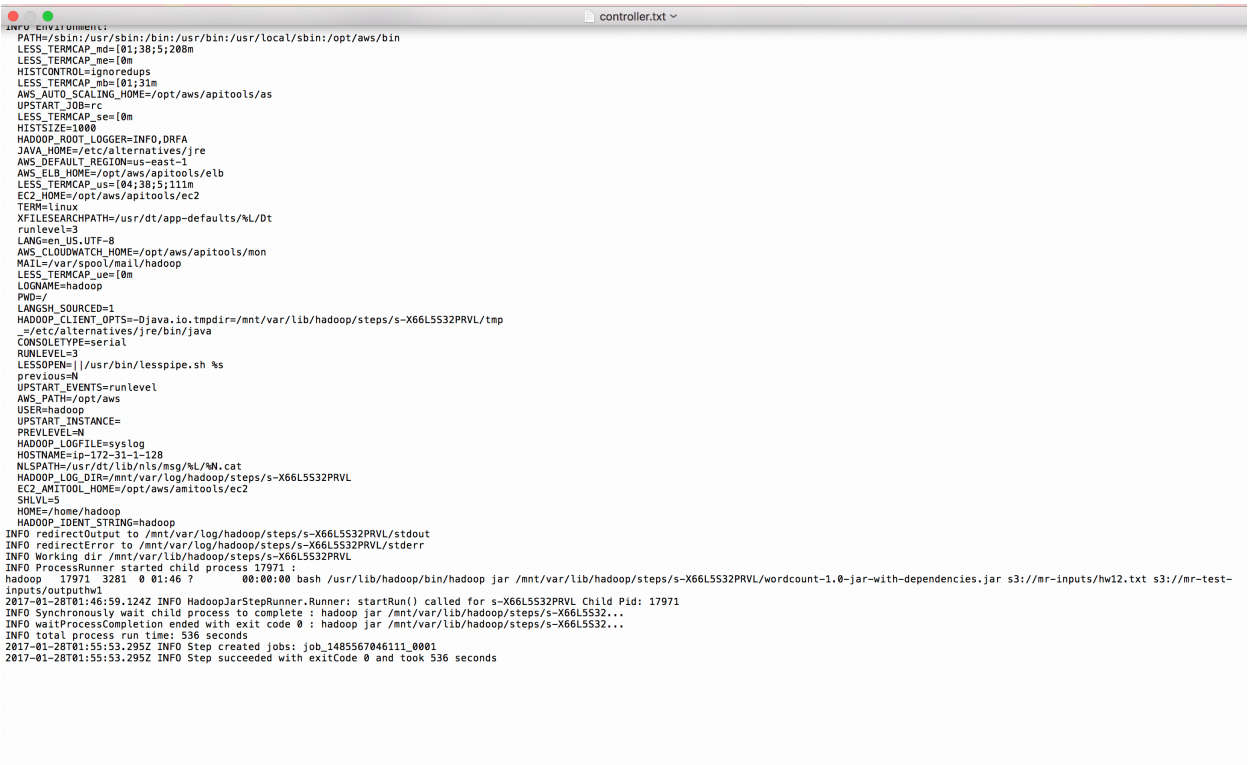
The modal also lists the jobs for the cluster s-X66L5S32PRVL:

Job	State	Start time (UTC-5)	Actions
job_1485567046111_0001	COMPLETED	2017-01-27 20:47 (UTC-5)	<a href="#">View tasks</a>

# Screenshot showing cluster details



# CONTROLLER.txt



# SYSLOG.txt

```
2017-01-28 01:55:40.433 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 96%
2017-01-28 01:55:43.444 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 97%
2017-01-28 01:55:45.444 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 98%
2017-01-28 01:55:49.452 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 99%
2017-01-28 01:55:51.457 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2017-01-28 01:55:51.460 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1485567846111_0001 completed successfully
2017-01-28 01:55:51.583 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 55

File System Counters
  FILE: Number of bytes read=291485730
  FILE: Number of bytes written=435847916
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=1936
  HDFS: Number of bytes written=0
  HDFS: Number of read operations=22
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=0
  S3: Number of bytes read=1454291074
  S3: Number of bytes written=72825
  S3: Number of read operations=0
  S3: Number of large read operations=0
  S3: Number of write operations=0

Job Counters
  Killed map tasks=1
  Launched map tasks=22
  Launched reduce tasks=3
  Data-local map tasks=22
  Total time spent by all maps in occupied slots (ms)=88651440
  Total time spent by all reduces in occupied slots (ms)=40327552
  Total time spent by all map tasks (ms)=1846905
  Total time spent by all reduce tasks (ms)=503412
  Total vcore-milliseconds taken by all map tasks=1846905
  Total vcore-milliseconds taken by all reduce tasks=503412
  Total megabyte-milliseconds taken by all map tasks=2836846080
  Total megabyte-milliseconds taken by all reduce tasks=1546481664

Map-Reduce Framework
  Map input records=21907700
  Map output records=262898800
  Map output bytes=2484011200
  Map output materialized bytes=143959744
  Input split bytes=1936
  Combine input records=0
  Combine output records=0
  Reduce input groups=5274
  Reduce shuffle bytes=143959744
  Reduce input records=262098800
  Reduce output records=5274
  Spilled Records=706206400
  Shuffled Maps =66
  Failed Shuffles=0
  Merged Map outputs=66
  GC time elapsed (ms)=19209
  CPU time spent (ms)=1102070
  Physical memory (bytes) snapshot=21772201984
  Virtual memory (bytes) snapshot=86583537664
  Total committed heap usage (bytes)=21357920256

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=1454291074
File Output Format Counters
  Bytes Written=72825
```